

Thermodynamic Sovereignty: A Reciprocal Resource Framework for Sovereign AI Infrastructure

1. Abstract

The physical expansion of artificial intelligence infrastructure has precipitated an acute geopolitical and environmental friction between raw compute capability and regional resource sovereignty. Traditional cloud deployments rely on an extractivist paradigm, consuming municipal resources while exporting algorithmic value to foreign monopolies. This paper introduces a socio-economic governing equation, $S = (C \cdot R_{\text{local}}) - E_{\text{entropy}} + K$, to formalize sovereign AI as a thermodynamic and resource-balancing problem. Within this framework, true sovereignty (S) is defined as a function of raw compute capacity (C) multiplied by a Local Resource Reciprocity Ratio (R_{local}), which measures the domestic retention of economic and operational value. This productive term is offset by Entropic Expansion Costs (E_{entropy}), representing localized negative externalities including grid degradation, water depletion, land displacement, and localized inflation. These dynamics are anchored by the Sovereign Constant (K), representing the institutional baseline of public policy, talent retention, and independent data residency laws. By analyzing global integrations—such as thermal energy networks in Europe and Canada alongside public compute initiatives in the Global South—this paper demonstrates that sustainable digital self-determination requires a shift from extractive "compute colonialism" to structured thermodynamic reciprocity. Policymakers must transition from passive infrastructure hosting to active systems orchestration, mandating high reciprocity coefficients and robust policy constants during development and zoning phases to achieve long-term macroeconomic and ecological stability.

2. Introduction and the Rise of Sovereign AI

The global computing landscape is undergoing a profound structural fragmentation, shifting away from centralized, globalized hyperscale cloud infrastructures toward highly localized, sovereign AI data clusters. Historically, the digital economy operated on an asymmetrical architecture where a small cohort of technology conglomerates based in the United States and China monopolized the physical and algorithmic layers of computation. This centralization has introduced severe geopolitical and macroeconomic vulnerabilities, exposing non-aligned states to the systemic threat of compute colonialism. Under this extractivist paradigm, transnational corporations extract raw, localized data and human labor from peripheral regions, process this material in foreign jurisdictions, and lease the resulting proprietary intelligence back to the host countries. This process depletes municipal utilities while reinforcing patterns of technological dependency and intellectual capital drain.

In response, national and municipal planners are increasingly asserting compute autonomy. This counter-movement is driven by the realization that next-generation economic competitiveness, national security, and epistemic self-determination are fundamentally contingent upon the domestic control of physical processing assets. For instance, algorithmic

colonialism frequently distorts local realities and devalues regional languages due to the dominance of English-centric training datasets. By establishing domestic infrastructure, nations seek to align foundation models with their specific linguistic, cultural, and constitutional values. This is illustrated by pan-African linguistic initiatives like the Masakhane project, which embeds *Ubuntu*-centered community values into localized language models, and Indian enterprises like Mokx, which integrate Vedas-inspired reasoning frameworks into local artificial intelligence systems.

However, the pursuit of digital self-reliance reveals an operational contradiction known as the Sovereign AI Paradox. While governments invest heavily in physical data centers to circumvent foreign cloud platforms, the underlying technology stack remains deeply reliant on global monopolies.

Layer of the Technology Stack	Ideal Sovereign Capability	Empirical Global Reality (2026 Data)	Systemic Vulnerability / Dependency
Hardware & Accelerators	Domestically fabricated silicon, advanced domestic packaging, and indigenous semiconductor intellectual property.	Complete dependency on a micro-monopoly; NVIDIA alone supplies GPUs for 52% of tracked sovereign projects.	Silicon supply-chain bottlenecks compromise absolute national self-determination.
Infrastructure Integration	Fully localized and state-owned data center operations independent of foreign cloud ecosystems.	Severe reliance on foreign technology providers; over 70% of sovereign projects involve at least one foreign partner.	Digital sovereignty is compromised at the physical and integration layer.
Model Architectures	Domestically trained frontier foundation models utilizing culturally and linguistically aligned datasets.	High usage of foreign open-weight architectures; Meta's Llama family is utilized in 36% of sovereign model projects.	Models risk inheriting foreign ethical biases, epistemologies, and security backdoors.
Capital Investment	Diverse, public-private domestic funding mechanisms distributed across regional hubs.	Extreme concentration of capital; the UAE and Japan account for over two-thirds of total disclosed global sovereign AI investments.	Smaller economies face a severe digital divide, unable to match the multi-billion dollar capital pools of leading nations.

This structural mapping demonstrates that absolute technological isolation is virtually impossible in the near term. Consequently, contemporary statecraft must conceptualize digital self-determination not as complete autarky, but as a layered sovereignty. In this framework, nations strategically select which layers of the technology stack to host domestically, while actively implementing regulatory and physical safeguards to mitigate foreign exploitation.

3. Deconstructing the Sovereign AI Equation

To evaluate the net socio-economic and thermodynamic yield of physical computing

infrastructure, this paper proposes a formalized governing equation:

This mathematical abstraction frames sovereignty as a dynamic resource-balancing problem, where positive outputs must exceed the physical and economic costs of computation.

The variable S represents Sovereignty, defined as the net level of localized self-determination, data control, cultural alignment, and strategic security achieved by a territory. Crucially, S is modeled as a thermodynamic yield rather than a static state. If the physical and economic drag of a computing facility exceeds its localized return, the net value of S becomes negative, indicating that the hosting territory has devolved into an exploited digital enclave. True sovereignty requires a positive net yield, demonstrating that the infrastructure actively enhances local civic and economic capabilities without depleting the physical environment.

The variable C represents Compute, the physical capacity scalar of the system. This term aggregates raw processing power, silicon infrastructure, and networking throughput capacity. While C acts as a primary engine of digital acceleration, it is fundamentally a physical multiplier. In the absence of localization mechanisms, a higher value of C simply increases the rate of resource consumption, multiplying the negative externalities imposed on the local community without guaranteeing any domestic benefit.

The variable R_{local} represents the Local Resource Reciprocity Ratio, which serves as the efficiency coefficient of the system. This term measures the proportion of model utility, economic value, and compute equity that is funneled back into the immediate municipal community, as opposed to being exported to global monopolies. When R_{local} approaches zero, the physical compute cluster operates as a purely extractive mechanism, drawing on local power and water while exporting its intelligence surplus. Conversely, when R_{local} is maximized, the compute infrastructure operates symbiotically, returning equivalent or greater thermodynamic, civic, and economic value to the hosting region.

The variable E_{entropy} represents the Entropic Expansion Costs, quantifying the inescapable physical and social drag generated by the computing infrastructure. This variable aggregates the localized negative externalities, including grid degradation, aquifer depletion, carbon emissions, public health damage, and specialized real estate inflation. It represents the thermodynamic tax of computing, representing the physical limits that raw silicon infrastructure imposes on regional ecosystems.

The variable K represents the Sovereign Constant, the non-physical, institutional baseline of the state or municipality. This variable encapsulates the efficacy of public policy, domestic talent retention, and independent data residency laws. Even when physical compute (C) is modest, a robust value of K ensures that the domestic ecosystem remains resilient against external technological coercion and brain drain, serving as the regulatory shield that prevents predatory resource extraction.

4. The Macroeconomic Theory of Surplus "Bleed"

The siting of modern computing facilities generates massive physical and economic spillovers that directly impact neighboring municipal and national communities. These spillovers can be mathematically modeled as a divergence between the positive multiplier effect of $(C \cdot R_{\text{local}})$ and the negative thermodynamic drag of E_{entropy} .

Positive Externalities and the Synergistic Spillover

When a municipality implements policies that mandate a high R_{local} ratio, the data center is

transformed from an isolated, resource-intensive enclave into a circular engine of industrial and civic ecology. The primary mechanism of this synergistic spillover is thermal reciprocity, facilitated by Thermal Energy Networks (TENs) and district heating systems. AI data centers convert nearly 100% of their electrical input into low-temperature waste heat, which has historically been rejected into the atmosphere via water-intensive evaporative cooling systems. By transitioning to closed-loop, direct-to-chip liquid cooling systems, facilities can capture high-grade thermal energy (ranging from 122°F to 140°F) and route it through underground, water-filled pipes directly into municipal district heating grids.

This thermal integration has been empirically validated across several European and North American municipalities. In Sweden, the Stockholm district heating scheme recovers server waste heat to warm 30,000 residential apartments annually, halving carbon emissions compared to traditional natural gas heating. Similarly, the Nebius Group data center in Mäntsälä, Finland, recovers enough thermal energy to heat 2,500 local homes , while the Tallaght District Heating Scheme in Ireland saved over 1,100 metric tonnes of carbon dioxide in 2023 by repurposing server heat for civic buildings. During the 2024 Summer Olympics in Paris, waste heat from a local data center was captured to heat the training pool at the Olympics Aquatics Center, with plans to scale the system to warm 1,000 homes. In British Columbia, Canada, the TELUS sovereign AI cluster redirects its thermal output to heat 150,000 homes in Vancouver.

Beyond residential heating, this thermal energy can be cascaded into agricultural and industrial co-locations. In West Virginia and Ohio, proposed industrial campuses utilize server waste heat and captured carbon dioxide to heat commercial greenhouses, providing a year-round, low-carbon heat supply that protects local food production from fossil fuel price spikes.

Furthermore, siphoning this low-temperature heat to assist anaerobic digestion in municipal wastewater treatment facilities reduces the auxiliary energy required to run regional sanitation systems. This shift from evaporative cooling to closed-loop heat recovery reduces a data center's water consumption by up to 80% and lowers its operational energy footprint by 10% to 30%, which at a national scale could prevent the construction of 54 new power plants and save utility customers \$22.1 billion in construction costs. This integration is particularly valuable for post-industrial transitions, such as converting legacy brownfields and shuttered coal plants into modern, co-located digital and manufacturing hubs.

Municipal Integration Method	Empirical Precedent / Location	Thermodynamic & Environmental Impact	Economic & Civic Benefit
District Heating Networks (TENs)	Stockholm, Sweden; Tallaght, Ireland; Vancouver, Canada.	Captures and redirects high-grade thermal energy, reducing municipal reliance on fossil-fuel heating.	Lowens consumer heating costs; reduces greenhouse gas emissions.
Agricultural Co-location	West Virginia Greenhouse Proposals; Ohio Agritech.	Utilizes waste heat and captured carbon dioxide for year-round greenhouse heating.	Stabilizes regional food supply chains; reduces agricultural operating costs.
Compute Franchising Programs	IndiaAI Mission; Canada's AI Compute Access Fund.	Allocates physical GPU capacity to regional academic and startup networks.	Prevents domestic brain drain; lowers the capital barrier for local digital innovation.
Green-Grid Additionality	TELUS Sovereign AI Factory (98% clean	Requires developer-financed	Expands regional grid capacity; lowers the

Municipal Integration Method	Empirical Precedent / Location	Thermodynamic & Environmental Impact	Economic & Civic Benefit
	energy integration).	investment in new carbon-free generation assets.	carbon intensity of the municipal grid.

Simultaneously, a high R_{local} ratio manifests through compute franchising, a mechanism that democratizes access to physical processing power. In a standard extractive hosting arrangement, 100% of a data center's capacity is leased to multinational enterprises, forcing local startups and researchers to burn through expensive cloud credits and rely on foreign-headquartered systems. Compute franchising legally reserves a fixed quota of physical GPU capacity for regional academic networks, public sector developers, and domestic small- and medium-sized enterprises. This is illustrated by the IndiaAI mission, which deployed a national pool of over 34,000 publicly funded GPUs to subsidize access for regional developers , and the United Kingdom's AI Research Resource (AIRR), which integrates domestic startups into national computing roadmaps. This localized access ensures that homegrown technological talent remains within the national borders, directly boosting the long-term Sovereign Constant (K).

Negative Externalities and the Entropic Drag

When computing infrastructure is deployed without strict local reciprocity mandates, the hosting community is subjected to severe physical and economic degradation, representing an unconstrained entropic drag (E_{entropy}). The physical expansion of hyperscale data centers has begun to destabilize local electrical grids. In the United States, data center electricity consumption grew from 76 TWh in 2018 to 176 TWh in 2023, representing approximately 4.4% of total national consumption. By 2028, the Lawrence Berkeley National Laboratory projects that this demand will rise to between 325 TWh and 580 TWh, consuming up to 12% of the national grid.

This rapid expansion has already triggered severe grid congestion. Up to 90 gigawatts of new data center capacity is projected to come online globally by 2030, introducing massive, continuous baseload demands that strain transmission infrastructure and force utilities to delay critical interconnection projects. The capital costs of these massive grid upgrades are routinely passed on to residential utility customers in the form of rate hikes, meaning local populations subsidize the physical infrastructure of transnational tech corporations. Furthermore, high-frequency power oscillations during large-scale training runs introduce severe reliability risks. In July 2024, a brief voltage fluctuation in northern Virginia triggered the simultaneous disconnection of 60 data centers, causing an instantaneous 1,500-megawatt power surplus that forced emergency grid adjustments to prevent a cascading regional blackout.

Local Resource Metric	Baseline Value (2021-2023)	Projected Horizon (2028-2030)	Systemic Socio-Economic & Environmental Cost
Grid Electricity Demand (US)	176 Terawatt-hours (4.4% of national grid).	325 to 580 Terawatt-hours (up to 12.0% of national grid).	Shifts infrastructure costs to residential ratepayers; delays grid interconnections.
Potable Water Consumption	Mid-sized sites consume 300,000	AI-related data centers projected to consume	Depletes regional aquifers; threatens

Local Resource Metric	Baseline Value (2021-2023)	Projected Horizon (2028-2030)	Systemic Socio-Economic & Environmental Cost
	gallons/day; large sites consume up to 5,000,000 gallons/day.	32 billion gallons annually.	municipal water supplies in water-stressed basins.
Land-Use Footprint	Average site covered 92 acres in 2022.	Average site covers 224 acres in 2024 (a 144% expansion).	Displaces regional agricultural and development land; inflates local real estate values.
Backup Generator Emissions	Negligible under standard grid conditions.	Emergent operations release high particulate matter and NO _x (200x-600x higher than gas plants).	Drives up to \$99 million in annual health damages and causes 3.4 to 6.5 premature deaths per site.

Water depletion represents another severe entropic cost. A typical mid-sized data center utilizing evaporative cooling towers withdraws approximately 300,000 gallons of water daily, while massive facilities consume up to 5,000,000 gallons daily—comparable to the consumption of a small town. By 2028, AI-related data centers in the United States are projected to consume 32 billion gallons of water annually, equivalent to the indoor water use of 360,000 households. This extraction is highly critical given that two-thirds of all data centers built or under development since 2022 are located in water-stressed zones, such as the Colorado River Basin, Texas, and southern Arizona. Corporate environmental reports confirm the magnitude of this extraction: in 2021, Google's data centers withdrew 4.2 billion gallons of water, consuming 3.3 billion gallons, while Meta's facilities withdrew 1.3 billion gallons, drawing heavily on highly stressed municipal aquifers.

This physical extraction is accompanied by substantial public health risks. To guarantee uninterrupted uptime, data centers maintain massive arrays of diesel-fueled backup generators. While only utilized during grid emergencies and testing, these generators emit highly concentrated fine particulate matter (PM_{2.5}) and nitrogen oxides (NO_x), releasing 200 to 600 times more NO_x than standard natural gas plants. Air pollution exposure models indicate that the operation of backup generators at a single massive facility can generate between \$53 million and \$99 million in annual localized health damages, driven by an estimated 3.4 to 6.5 additional premature deaths per year alongside increased hospitalizations and lost productivity. At a regional scale, even limited emergency generator use in northern Virginia is associated with nearly \$300 million in annual public health costs across multiple states.

Finally, unconstrained entropic expansion drives tax base fragmentation and local economic asymmetry. To attract technology developers, state and municipal authorities frequently grant long-term property and sales tax abatements. However, data centers are highly capital-intensive but fundamentally non-labor-intensive, generating very few permanent, high-quality jobs (averaging just 27 to 111 permanent employees per facility). Because the massive corporate profits generated by the infrastructure are immediately exported to global metropolises, host communities are left to shoulder the long-term utility, infrastructure, and environmental costs, while regional schools and public services are deprived of vital local tax revenues.

5. Policy Framework and Real-World Applications

To counter the threat of neocolonial digital extraction, national and municipal planners must

actively manipulate the variables of the sovereign equation. This requires implementing concrete policy interventions designed to artificially maximize the Sovereign Constant (K) and legally mandate high Local Resource Reciprocity Ratios (R_{local}) during zoning, permitting, and development phases.

Artificially Boosting K (Institutional Hardening)

Strengthening the Sovereign Constant requires establishing a robust legal, educational, and operational framework that protects domestic resources from foreign capture. The primary mechanism for institutional hardening is the codification of operational network controls. Rather than merely permitting foreign technology entities to build and manage regional infrastructure, governments should adopt policies modeled on Canada's Sovereign AI Compute Strategy and the "Enabling Large-Scale Sovereign AI Data Centres" initiative. Under these regulatory guidelines, any data center facility exceeding a total capacity of 100 megawatts must satisfy strict sovereignty factors, including demonstrating direct indigenous and local participation, utilizing domestic partners, and ensuring that the physical operations and network security are controlled by domestic firms under domestic legal jurisdiction.

This institutional model recognizes that digital sovereignty does not mean complete isolation; rather, it allows host nations to leverage world-class technology—such as importing advanced graphics processors—while legally retaining control of data residency and governance. Furthermore, policymakers must address the digital rights and regulatory gaps that leave emerging economies vulnerable to data exploitation. In sub-Saharan Africa, national frameworks like Kenya's Data Protection Act represent critical state efforts to assert control over domestic digital ecosystems. However, the practical effectiveness of these frameworks is heavily constrained by a lack of specialized technical expertise within civil society organizations, funding precarity, and intense political pressure from foreign-funded smart city and biometric surveillance procurement programs.

To overcome these structural limitations, governments must institutionalize public participation and invest in regional regulatory bodies, such as the proposed Pan-African AI Sovereignty Agency. By pooling technical resources, developing regional cultural data libraries, and establishing joint-venture supercomputing platforms, sovereign coalitions can negotiate equitable terms with global technology providers and resist the imposition of foreign-dominated standard-setting frameworks.

Mandating High R_{local} (Thermodynamic Permitting & Siting)

To transition computing facilities from resource drains into reciprocal civic assets, municipal zoning boards must codify thermodynamic permitting guidelines. First, zoning permits must be legally conditioned on the installation of waste-heat recovery systems and integration into a local district heating network. This can be achieved by adopting regulatory precedents like the German Energy Efficiency Act, which mandates strict energy audits and waste-heat reuse for all new data centers, or New York's S9422 (the *Utility Thermal Energy Network and Jobs Act*), which directs public utilities to construct and manage neighborhood-scale thermal energy networks. Under these frameworks, developers are restricted from using open-evaporative cooling towers, forcing a transition to closed-loop liquid cooling systems that minimize local aquifer strain while providing a high-grade thermal supply for nearby buildings, greenhouses, or manufacturing facilities.

Second, municipal planners must enforce strict grid additionality clauses. Data center

developers must be legally prohibited from drawing on existing grid capacity, which can cause consumer utility rate hikes and prolong the operation of fossil-fuel-burning power plants. Permits must require developers to finance and build new, localized carbon-free energy assets (such as solar, wind, or geothermal generation) equivalent to 100% of their operational load, while integrating behind-the-meter battery storage to support peak-load shedding.

Third, municipalities must implement compute franchising allocations as a core condition of site authorization. A fixed percentage of the facility's active processing capacity must be routed to a "common compute pool," managed by local public institutions and made accessible to regional universities, public schools, and local tech startups at subsidized or "at-cost" rates. This directly mitigates the financial precarity of regional founders, who routinely exhaust their limited budgets on expensive foreign cloud services.

Policy Instrument	Precedent / Statutory Model	Target Variable	Operational Mechanism
Mandatory Waste-Heat Recovery	German Energy Efficiency Act.	Maximizes $R_{\{local\}}$	Legally binds developers to capture server heat and route it to local municipal or industrial networks.
Utility Thermal Network Mandates	New York S9422 (<i>Utility Thermal Energy Network and Jobs Act</i>).	Minimizes $E_{\{entropy\}}$	Directs regulated utilities to build and manage neighborhood-scale underground thermal networks.
Operational Control Requirements	Canadian Sovereign AI Compute Strategy (ISED).	Artificially boosts K	Restricts permits to projects owned and operated by domestic firms under domestic jurisdiction.
Grid Additionality Clauses	Standard Western US Grid Zoning Ordinances.	Minimizes $E_{\{entropy\}}$	Prohibits data centers from drawing existing capacity; forces investment in new renewable generation.
Compute Franchising Mandates	India AI Mission GPU Common Pool.	Maximizes $R_{\{local\}}$	Reserves a fixed quota of physical GPU processing capacity for local academic and startup networks.

6. Conclusion

The conceptualization of "Sovereign AI" is fundamentally a physical, thermodynamic, and macroeconomic problem rather than a mere software or algorithmic challenge. The explosive expansion of physical computing assets is placing unprecedented, continuous stress on municipal electrical grids, regional water tables, and public health. When managed under traditional extractivist frameworks, uncoordinated compute expansion functions as a powerful engine of entropic drag ($E_{\{entropy\}}$), draining local resources while exporting the resulting

high-value intelligence and economic surplus to a small number of global monopolies. True digital sovereignty (S) cannot be achieved through isolation or physical hosting alone. Instead, it requires the rigorous implementation of policy and structural frameworks that mandate a high Local Resource Reciprocity Ratio (R_{local}) while strengthening the Sovereign Constant (K). By treating server waste heat as a valuable public utility, demanding strict grid additionality, and franchising processing capacity to regional innovators, national and municipal planners can balance the thermodynamic and economic scales of digital expansion. Ultimately, data centers must be integrated into the circular industrial ecology of their host communities, transforming these facilities from resource enclaves into reciprocal municipal assets that anchor clean, resilient, and autonomous regional economic systems.

Works cited

1. The Illusion of AI Sovereignty: Washington and Beijing Still Pull the Strings - Avasant, <https://avasant.com/report/the-illusion-of-ai-sovereignty-washington-and-beijing-still-pull-the-strings/>
2. Sovereign AI: Building ecosystems for strategic resilience and impact - McKinsey, <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/sovereign-ai-building-ecosystems-for-strategic-resilience-and-impact>
3. How the Global South is reimagining the future of AI - The World Economic Forum, <https://www.weforum.org/stories/2026/02/how-the-global-south-is-reimagining-the-future-of-ai/>
4. Algorithmic Colonialism | Politikon: The IAPSS Journal of Political Science, <https://politikon.iapss.org/index.php/politikon/article/view/526>
5. State of Compute Access 2024: How to Navigate the New Power Paradox, <https://institute.global/insights/tech-and-digitalisation/state-of-compute-access-2024-how-to-navigate-the-new-power-paradox>
6. The hidden cost of AI: Digital colonialism and the Global South, <https://waccglobal.org/the-hidden-cost-of-ai-digital-colonialism-and-the-global-south/>
7. DIGITAL SOVEREIGNTY AND DATA COLONIALISM: SHAPING A JUST DIGITAL ORDER FOR THE GLOBAL SOUTH - Policy Center, https://www.policycenter.ma/sites/default/files/2025-10/PP_38-25%20%28Marcus%20Vini%CC%81cius%20De%20Freitas%29.pdf
8. Algorithmic dependence and digital colonialism: a conceptual framework for artificial intelligence in education and knowledge systems of the Global South - Frontiers, <https://www.frontiersin.org/journals/education/articles/10.3389/feduc.2026.1720563/full>
9. Digital Colonialism and the Role of Local Intermediaries: Examining Big Tech's Impact on Data Sovereignty and Human Rights in Africa, <https://www.cambridge.org/core/journals/business-and-human-rights-journal/article/digital-colonialism-and-the-role-of-local-intermediaries-examining-big-techs-impact-on-data-sovereignty-and-human-rights-in-africa/635F708D2C0E9D5832593A7366BCC6C9>
10. UK Compute Roadmap - GOV.UK, <https://www.gov.uk/government/publications/uk-compute-roadmap/uk-compute-roadmap>
11. Sovereign AI Index - CNAS Reports, <https://interactives.cnas.org/reports/sovereign-ai-index/>
12. Canada wants sovereign AI data centres. What does that actually mean? | CBC News, <https://www.cbc.ca/news/business/bakx-canada-sovereign-ai-9.7198649>
13. Using Waste Heat Generated from Data Centers: - ReImagine ..., https://reimagineappalachia.org/wp-content/uploads/2026/02/Catching-Heat_Using-Waste-Heat-Generated-from-Data-Centers.pdf
14. AI, Data Centers, and the U.S. Electric Grid: A Watershed Moment, <https://www.belfercenter.org/research-analysis/ai-data-centers-us-electric-grid>
15. WHAT HAPPENS WHEN DATA CENTERS COME TO TOWN? - Science, Technology, and

Public Policy,

<https://stpp.fordschool.umich.edu/sites/stpp/files/2025-07/stpp-data-centers-2025.pdf> 16.

Thermal Energy Networks Turn Data Center Waste Heat into a Hot ...,

<https://www.eesi.org/articles/view/thermal-energy-networks-turn-data-center-waste-heat-into-a-hot-commodity> 17. Hard to Compute - Why Startups Need More Power,

<https://api.startupcoalition.io/u/2025/11/Hard-to-Compute-Why-Startups-Need-More-Power-2.pdf>

18. From Energy Use to Air Quality, the Many Ways Data Centers Affect US Communities,

<https://www.wri.org/insights/us-data-center-growth-impacts> 19. Analyzing air pollution health,

economic risks from AI data centers,

<https://hsph.harvard.edu/news/analyzing-air-pollution-health-economic-risks-from-ai-data-centers/> 20. AI Sovereign Compute Infrastructure Program - OECD.AI,

<https://oecd.ai/en/dashboards/policy-initiatives/ai-sovereign-compute-infrastructure-program> 21.

Resisting Digital Colonialism and Protecting Subaltern Agency in ...,

<https://jpia.princeton.edu/news/resisting-digital-colonialism-and-protecting-subaltern-agency-africa> 22. Sector Coupling of Data Centers & District Heating - nyserda,

<https://www.nyserda.ny.gov/-/media/Project/Nyserda/Files/Programs/Large-Scale-Thermal/DATA-HEAT---Sector-Coupling-Data-Centers-and-District-Energy.pdf> 23. TELUS' Clean Energy &

Water-Reduced AI Data Centre in Canada ...,

<https://sustainabilitymag.com/news/telus-clean-energy-water-reduced-ai-data-centre-in-canada>

24. GB cloud: building the UK's compute capacity - Centre for British Progress,

<https://britishprogress.org/uk-day-one/gb-cloud-building-the-uks-compute-capacity> 25. Are data centers depleting the Southwest's water and energy resources?,

<https://www.apmresearchlab.org/10x/data-centers-resource> 26. Enabling large-scale sovereign AI data centres - Government of Canada,

<https://ised-isde.canada.ca/site/ised/en/enabling-large-scale-sovereign-ai-data-centres>